

DistillPDF Demo Document

Abstract

This document is a self-contained demo used by the distillpdf test suite. It exercises titles, headings, paragraphs, lists, tables and figures so the extraction pipeline can be verified without any third-party PDFs.

1 Introduction

DistillPDF reconstructs the logical structure of a PDF and emits clean, semantic HTML. This introduction is plain body prose that should be extracted as ordinary paragraphs in reading order.

The following sections demonstrate a bulleted list, a data table and an embedded figure, each with known content the tests assert against.

- * Headings are recovered from font size and weight.
- * Tables are reconstructed into rows and cells.
- * Figures are wrapped together with their captions.

2 Methods

We render this document with a fixed layout so its extracted structure is deterministic and can serve as a stable regression fixture.

Metric	Value	Unit
Recall	0.98	ratio
Latency	12	ms




Figure 1: A simple demo chart with two bars.